IMPROVING OPTIMALITY AND EFFICIENCY OF GREEDY GROUP RECURSION ALGORITHM

Florin Dobrian [1] **Oleg Puzyrko White** [1]

ABSTRACT

Recently (Liu et al., 2025) has presented efficient algorithm - Greedy Group Recursion (GGR) - for reordering the rows and the fields within each row of an input table to maximize key-value (KV) cache reuse when performing LLM serving. In this paper, we propose several adjustments to GGR algorithm that can improve optimality of the solution and reduce its execution time.

1 INTRODUCTION

There has been growing research on LLM inference optimization. In particular, recent work (Liu et al., 2025; Cheng et al., 2025) presents solutions to optimize relational data analytics workloads for offline LLM inference. It proposes Greedy Group Recursion (GGR), an approximate algorithm that leverages functional dependencies (such as primary and foreign key relationships from the data schema) and table statistics, which are readily available in many databases and analytics systems, to reduce the search space.

1.1 Greedy Group Recursion (GGR) Algorithm

The GGR algorithm is described in detail in (Liu et al., 2025). Let us briefly outline its main points.

The pseudocode specification of GGR is in Algorithm 1 listing and is taken from the original paper but with several critical corrections (and multiple little typo fixes).

GGR algorithm optimizes LLM queries by finding a $(n \times m)$ table rearrangement that maximizes the LLM's KV cache prefix hit count (PHC). It achieves this goal by reordering the rows and the fields within each row. Each row may have a different field order. The rearranged table is represented as a list of tuples L, where each tuple in L corresponds to a row, and the tuple elements contain the column values. A cell in the list of tuples is denoted as $L[r][f]$, indicating the value in tuple r at position f.

1.1.1 Prefix Hit Count Calculation

The hit count of a single cell $L[r][c]$ is non-zero **only** if its value is the same as in the previous row $L[r][c] = L[r-1][c]$ **and** all preceding fields have the same property $\forall f \leq$

---

[1]Data Analytics Group. Corresponding author: Oleg P. White <oleg.p.white@gmail.com>.

$c, L[r][f] = L[r-1][f]$. Then its hit count is computed as the square of the value's string length:

$$hit(L, r, c) = \begin{cases} \text{len}(L[r][c])^2 & \text{if } \forall f \leq c, \\ & L[r][f] = \\ & L[r-1][f] \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

The squared string length reflects the quadratic complexity of token processing in LLM inference, where each token computation depends on every preceding token and increases computational cost quadratically with the input length.

The hit count for a single row r in L is then just a sum of hit counts for all cells in a row:

$$hit(L, r) = \sum_{c=1}^{m} hit(L, r, c) \quad (2)$$

The PHC for a list of tuples L with n rows and m fields is given by:

$$\text{PHC}(L) = \sum_{r=1}^{n} hit(L, r) \quad (3)$$

The $\text{PHC}(L)$ is the objective function of a table reordering algorithm: it tries to find output L with maximum $\text{PHC}(L)$ value.

1.1.2 Selecting the highest-hit distinct value

The greedy part of GGR is based on finding a distinct value v (in the corresponding column c) with the highest hit count among all distinct values in the table T. If $R_v = R(v, c, T)$ are the rows with value v in the column c:

$$R(v, c, T) = \{i \mid T[i, c] = v\}, \quad (4)$$

then hit count of v, c is given by:

$$HitCount(v, c, T) = \text{len}(v)^2 \times (|R(v, c, T)| - 1) \quad (5)$$

Algorithm 1 Greedy Group Recursion (GGR)

1: **Input:** Table T, Functional Dependency FD
2: **Output:** Prefix Hit Count S, Reordered List of Tuples L

3: **function** HITCOUNT(v, c, T, FD)
4:     $R_v \leftarrow \{i \mid T[i, c] = v\}$
5:     inferred_cols $\leftarrow \{c' \mid (c, c') \in FD\}$
6:     inferred_vals $\leftarrow \{T[R_v[1], c'] \mid c' \in \text{inferred\_cols}\}$
7:     tot_len $\leftarrow \text{len}(v)^2 + \sum_{v' \in \text{inferred\_vals}} \text{len}(v')^2$
8:     $HC \leftarrow \text{tot\_len} \times (|R_v| - 1)$
9:     $cols \leftarrow [c] + \text{inferred\_cols}$
10:    $vals \leftarrow [v] + \text{inferred\_vals}$
11:    **return** $HC, cols, vals$
12: **end function**

13: **function** GGR(T, FD)
14:    **if** $|T|_{rows} = 1$ **then**
15:        **return** $0, [T[1]]$
16:    **end if**
17:    **if** $|T|_{cols} = 1$ **then**
18:        $S \leftarrow \sum_{v \in \text{distinct}(T[,1])} \text{HITCOUNT}(v, 1, T)$
19:        **return** $S, sort([T[i] \mid i \in 1 \ldots |T|_{rows}])$
20:    **end if**
21:    $b\_v, b\_c \leftarrow \text{None}, \text{None}$
22:    $max\_HC, b\_cols, b\_vals \leftarrow -1, [], []$
23:    **for** $c \in \text{columns}(T), v \in \text{distinct}(T[, c])$ **do**
24:        $HC, cols, vals \leftarrow \text{HITCOUNT}(v, c, T, FD)$
25:        **if** $HC > max\_HC$ **then**
26:            $b\_v, b\_c \leftarrow v, c$
27:            $max\_HC, b\_cols, b\_vals \leftarrow HC, cols, vals$
28:        **end if**
29:    **end for**
30:    $R_v \leftarrow \{i \mid T[i, b\_c] = b\_v\}$
31:    $A\_HC, A\_L \leftarrow \text{GGR}(T[\text{rows} \setminus R_v, \text{cols}], FD)$
32:    $B\_HC, B\_L \leftarrow \text{GGR}(T[R_v, \text{cols} \setminus b\_cols], FD)$
33:    $C\_HC \leftarrow max\_HC$
34:    $S \leftarrow A\_HC + B\_HC + C\_HC$
35:    $L \leftarrow [b\_vals + B\_L[i] \mid i \in 1 \ldots |R_v|] + A\_L$
36:    **return** S, L
37: **end function**

38: **return** GGR(T, FD)

GGR then places all v, c cells in consecutive rows and in the first field of the output L and splits remaining cells of the table T into the sub-table of rows R_v with excluded column c: $T[R_v, \text{cols} \setminus [c]]$, and the sub-table of rows excluding R_v: $T[\text{rows} \setminus R_v, \text{cols}]$. GGR then recursively runs on each of two sub-tables. See Figure 1 for illustration of this key step.

Compared to the brute-force algorithm that requires $n! \times (m!)^n$ potential orderings, GGR significantly reduces the search space by selecting the highest-hit value and then reducing the dimensions of the table at each recursive step with the maximum depth of recursion $O(\min(n, m))$. GGR achieves close-to-perfect PHC output with fast practical execution time.
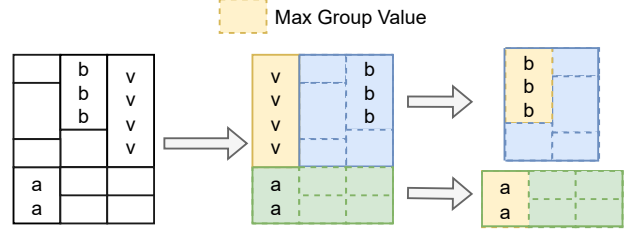

Figure 1. GGR picks the group with the maximum hit count at each step and calculates PHC as the sum of PHC of the elected group values (yellow box), the sub-table T excluding rows R_v (green box), and the sub-table of rows R_v excluding the field where the value is located in (blue box).

1.1.3 Functional Dependencies

GGR algorithm leverages a table's functional dependencies (FD) to reduce the number of fields it needs to consider at each recursion step.

We say that two columns A and B are in functional dependency $A \leftrightarrow B$ in a table T if, for any two rows r_1 and r_2, $T[r_1][A] = T[r_2][A]$ implies $T[r_1][B] = T[r_2][B]$ and vice versa.

In other words, for any distinct value a in the column A, there is a distinct value b in the column B such that they both are on the same rows. I.e., for $R_a \leftarrow \{i \mid T[i, A] = a\}$ and $R_b \leftarrow \{i \mid T[i, B] = b\}$ we have $R_a = R_b$.

FD rules can be specified as a list of disjoint sets containing column indices. For example, for the table with columns A, B, C, D, E, F where $A \leftrightarrow B$ and $C \leftrightarrow D \leftrightarrow E$, the FD rules can be written as $[[A, B], [C, D, E]]$.

GGR takes advantage of FD by combining distinct values from columns in the same FD rule into one block of group values.

1.1.4 GGR Specification

We can follow the pseudocode specification of GGR in Algorithm 1 listing.

At each recursive step, the GGR algorithm scans the table (lines 21–29) to find all distinct values with corresponding hit counts (lines 3–12). It then selects the highest-hit value b_v (in the column b_c) and splits the table into two sub-tables - one with all the rows R_v containing b_v value but excluding the column b_c and its FD-associated columns (line 31: $T[R\_v, \text{cols} \setminus b\_cols]$), and another sub-table with the remaining rows (line 32: $T[\text{rows} \setminus R\_v, \text{cols}]$). In the final step, GGR recurses on the two sub-tables (lines 31–32) and calculates the total PHC as the sum of PHCs computed for each sub-table and of hit count computed for b_v, b_c (line 34).

1.2 Reference Implementation of GGR Algorithm

The reference implementation of GGR specification from Algorithm 1 can be found in (ggr, 2026).

REFERENCES

GGR algorithm implementation in Python, 2026. URL https://github.com/whitechno/llm-sql-phc.

Cheng, A., Liu, S., Pan, M., Li, Z., Agarwal, S., Cemri, M., Wang, B., Krentsel, A., Xia, T., Park, J., Yang, S., Chen, J., Agrawal, L., Naren, A., Li, S., Ma, R., Desai, A., Xing, J., Sen, K., Zaharia, M., and Stoica, I. Let the barbarians in: How ai can accelerate systems performance research, 2025. URL https://arxiv.org/abs/2512.14806.

Liu, S., Biswal, A., Kamsetty, A., Cheng, A., Schroeder, L. G., Patel, L., Cao, S., Mo, X., Stoica, I., Gonzalez, J. E., and Zaharia, M. Optimizing llm queries in relational data analytics workloads, 2025. URL https://arxiv.org/abs/2403.05821.